

NN for sequential data cntd.

Beate Sick


sick@zhaw.ch

Remark: Much of the material has been developed together with Elvis Murina and Oliver Dürr

Topics

- **Recap architectures for sequential data**
 - 1D convolution
 - RNN
- **Recurrent NN with better memory**
 - GRU
 - LSTM

Task in homework: Predict how series will continue



Recap 1D “causal” convolution for ordered data

Toy example:

Output:


23	50	80
----	----	----

1D Kernel:

1	2
---	---

Input:

10	20	30
----	----	----




To make all layers the **same** size, a **zero padding** is added to the beginning of the input layers

“causal” networks, because the architecture ensured that no information from the future is used.

Stacking 1D “causal” convolutions without dilation


Non dilated Causal Convolutions



Stacking k causal 1D convolutions with kernel size 2 allows to look back k time-steps.
After 4 layers each neuron has a “memory” of 5 time-steps (1 present and 4 past).

Dilation allows to increase “memory” = receptive field

To increase the memory of neurons in the output layer, you can use “dilated” convolutions:





$$ReceptiveField = f_0 \cdot d_0 + \sum_i (f_i - 1) \cdot d_i = 2 \cdot 1 + 1 \cdot 2 + 1 \cdot 4 + 1 \cdot 8 = 16$$

Here the filter $f_i=2$ for all layer, but dilation is d_i starts with 1 and doubles from layer to layer

After 4 layers each neuron has a receptive field of 16 input neurons.

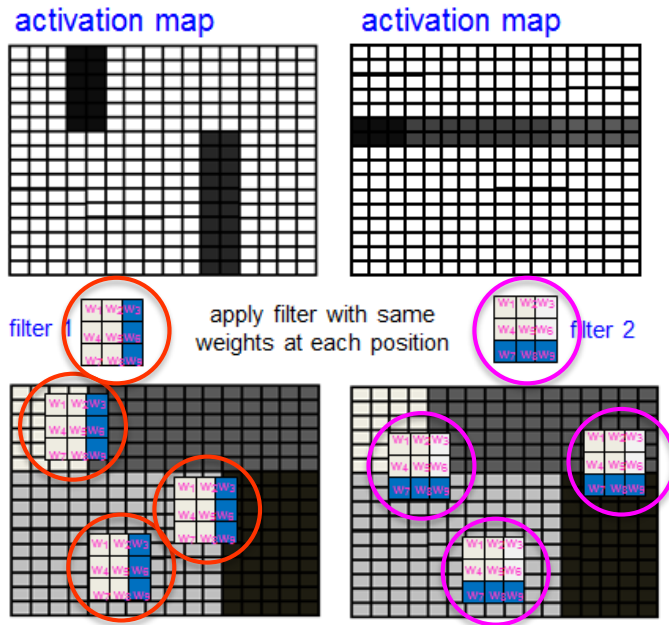
Recap the architecture of a simple RNN




$$\text{output}=\mathbf{h}_t = \tanh\left([\mathbf{h}_{t-1}, \mathbf{x}_t] \cdot \mathbf{W} + \mathbf{b}\right)$$

Common tricks in RNN & CNN and some differences

CNN and Recurrent Network share weights




CNN share weights between different local regions of the image



RNN share weights between time steps

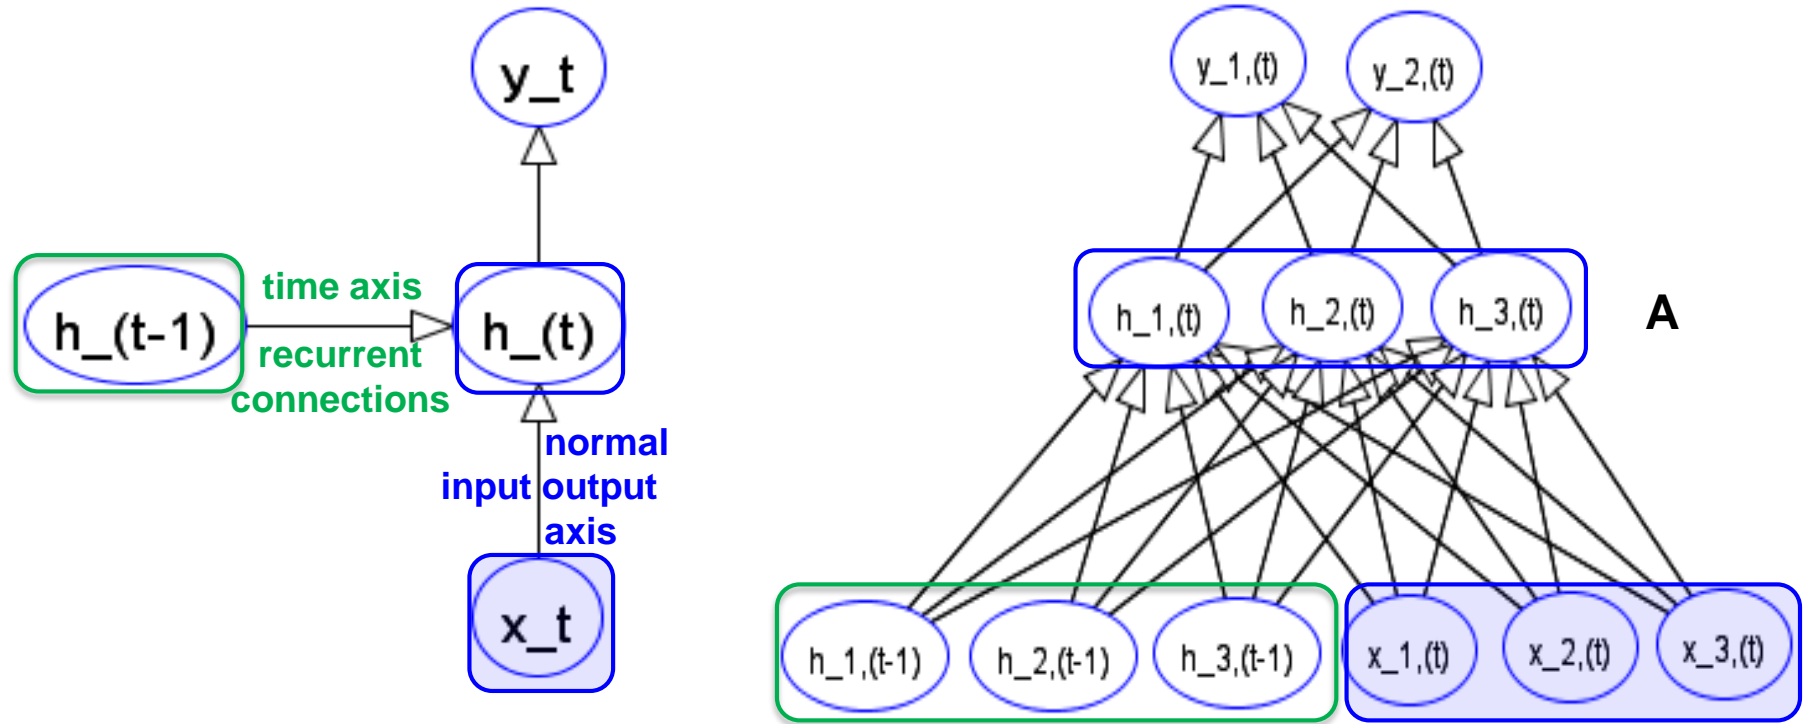
Remark: no weight sharing in fully connected NN

Also in RNN we can go deep for hierarchical features



Usually we see only 1-4 hidden layers in an RNN compared to usually 4-100 stacked hidden convolutional blocks in CNNs.

Dropout in recurrent architectures allow to choose different different dropout rates for recurrent and normal connections



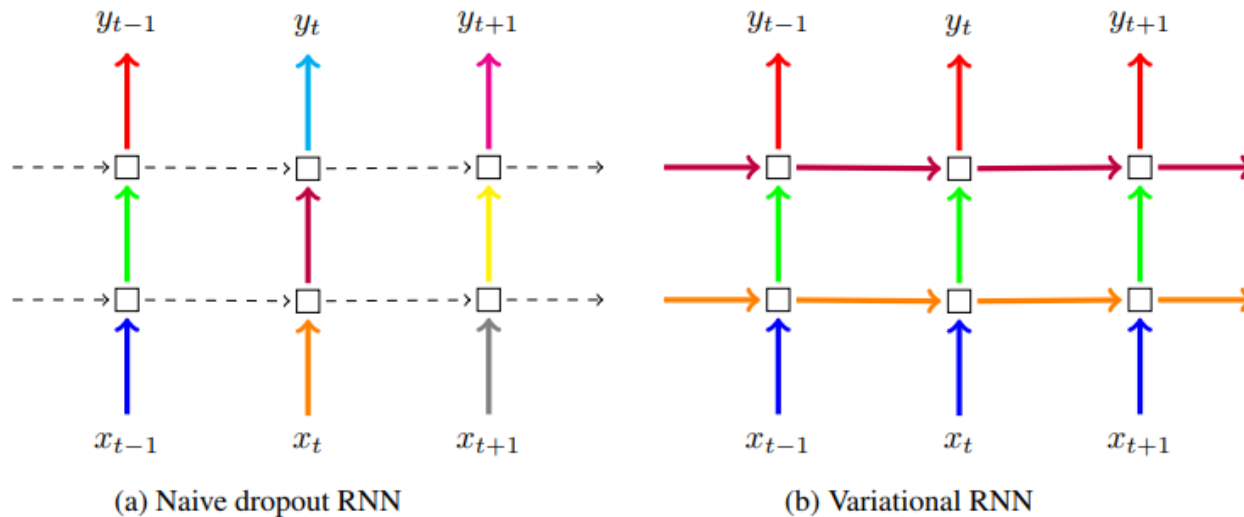
$$\mathbf{A} = f_{\mathbf{W}}(\mathbf{h}_{t-1}, \mathbf{x}_t) = \tanh([\mathbf{h}_{t-1}, \mathbf{x}_t] \cdot \mathbf{W} + \mathbf{b}) = \tanh(\mathbf{h}_{t-1} \cdot \mathbf{W}_h + \mathbf{x}_t \cdot \mathbf{W}_x + \mathbf{b})$$

$$\mathbf{W} = \begin{pmatrix} \mathbf{W}_h \\ \mathbf{W}_x \end{pmatrix}$$

Dimensions in example: \mathbf{W} :6x3, \mathbf{W}_h :3x3, \mathbf{W}_x :3x3

Dropout in recurrent architectures

It is important to **use identical dropout masks** (marked by arrows with same color) **at different time steps** in recurrent architectures like GRU or LSTM.



same arrow
color indicates
identical dropout
mask

Figure 1: **Depiction of the dropout technique following our Bayesian interpretation (right) compared to the standard technique in the field (left).** Each square represents an RNN unit, with horizontal arrows representing time dependence (recurrent connections). Vertical arrows represent the input and output to each RNN unit. Coloured connections represent dropped-out inputs, with different colours corresponding to different dropout masks. Dashed lines correspond to standard connections with no dropout. Current techniques (naive dropout, left) use different masks at different time steps, with no dropout on the recurrent layers. The proposed technique (Variational RNN, right) uses the same dropout mask at each time step, including the recurrent layers.

[Gal2016](#)

In keras:

```
model.add(layers.GRU(32, dropout=0.2, recurrent_dropout=0.2, input_shape=(None, ...)))
```

Dropout can fight overfitting in CNN and recurrent NN




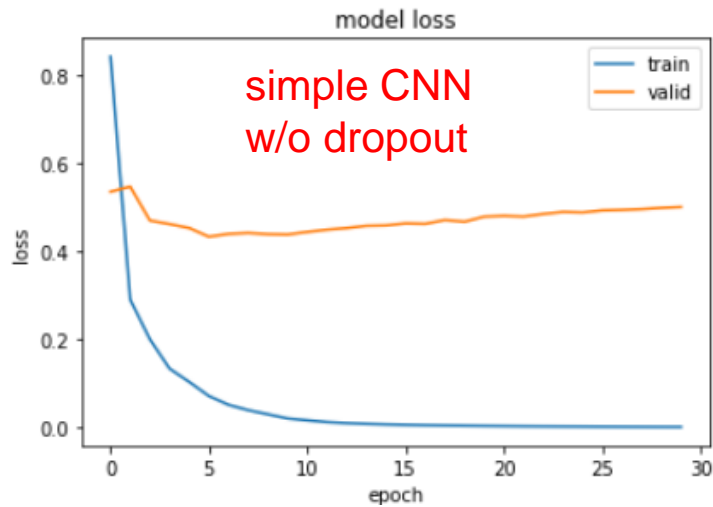
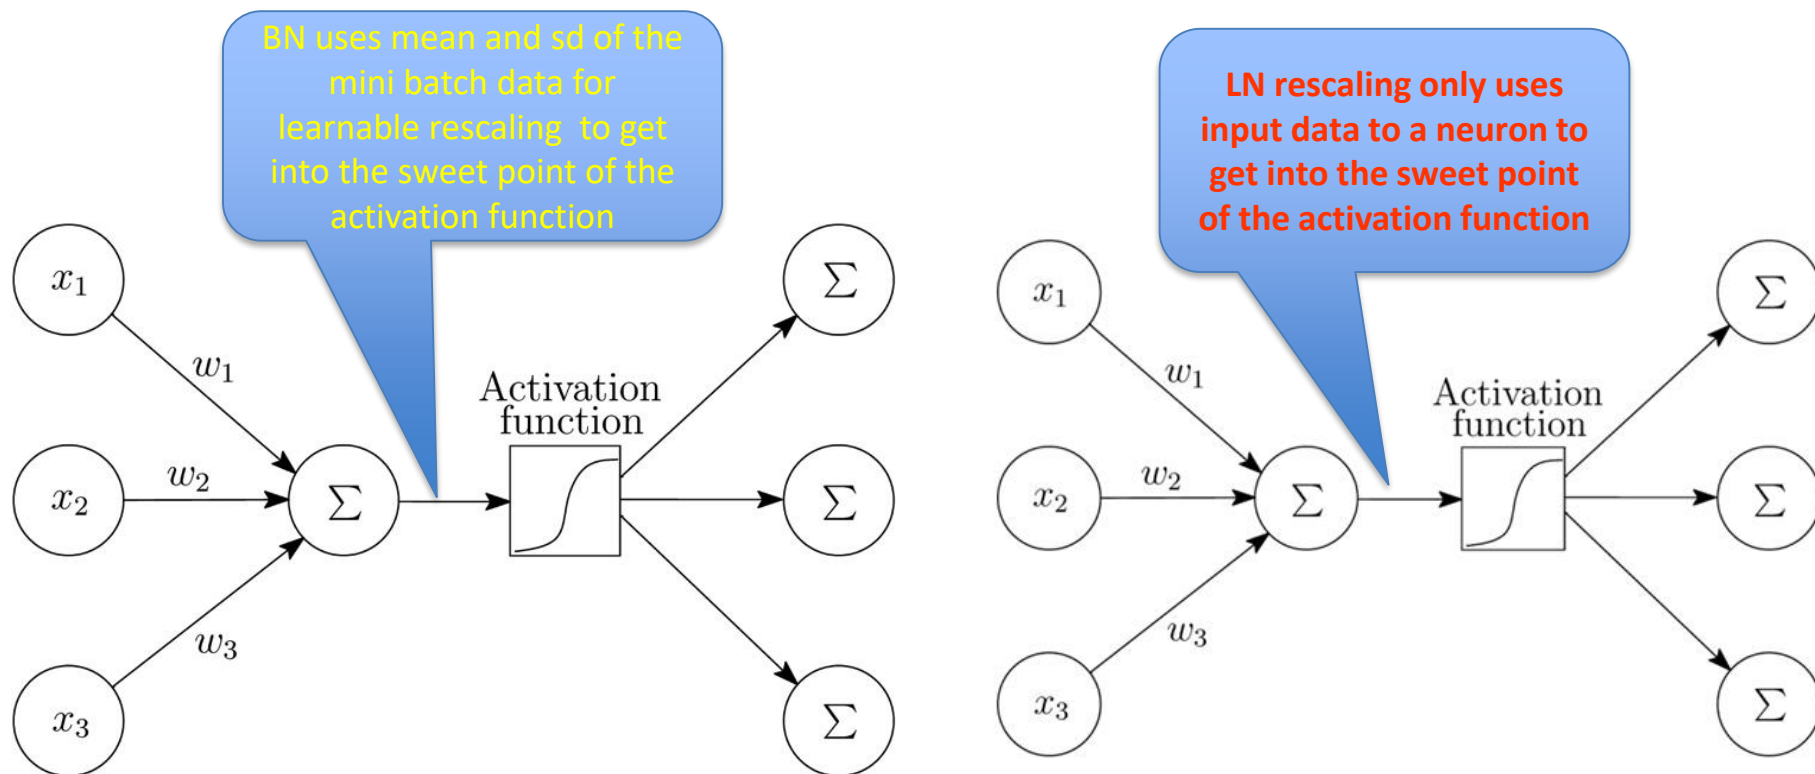


Image credits: F Chollet's book: DL with Python

Batchnormalization is crucial to train deep CNNs

Layernormalization is beneficial in RNN: LN \neq BN



Applying BN to RNN would not take into account the recurrent architecture of the NN over which statistics of the input to a neuron might change considerable within the same mini batch. In LN the mean and variance from all of the summed inputs to the neurons in a layer on a single training case are used for normalization .


Stateful RNN model

Training a stateful RNNs

- RNN are often trained on sequence data with inherent order
- Sequences are often very long and need to be cut between mini-batches
- By default the hidden state is initialized with zeros in each mini-batch
- In stateful RNN we connect sequences in the right order between mini-batches allowing to make use of the hidden state learned so far
- This requires a careful construction of the mini-batches and an appropriate transfer of the hidden state between mini-batches

Mini-batches in statefull RNN

The gradient is propagated back a fixed amount of steps defined by the size of a mini-batch. In stateful RNNs the hidden state is carried over between mini-batches and hence between connecting sequences given appropriate batches.



Vanishing/Exploding Gradient problem during training a RNN

Recall: Loss of a mini-batch is used to determine update

mini-batch of size M=8

train data input (S=len(seq)=3):

instance_id	seq_t1	seq_t2	seq_t3
1	\mathbf{x}_{11}	\mathbf{x}_{12}	\mathbf{x}_{13}
2	\mathbf{x}_{21}	\mathbf{x}_{22}	\mathbf{x}_{23}
3	\mathbf{x}_{31}	\mathbf{x}_{32}	\mathbf{x}_{33}
⋮	⋮	⋮	⋮
8	\mathbf{x}_{81}	\mathbf{x}_{82}	\mathbf{x}_{83}

train data target (2 classes, K=2):



instance_id	y_t1	y_t2	y_t3
1	(1,0)	(1,0)	(0,1)
2	(0,1)	(1,0)	(0,1)
3	(0,1)	(0,1)	-1
⋮	⋮	⋮	⋮
8	(1,0)	(1,0)	(1,0)

Cost C or Loss is given by the cross-entropy averaged over all instances in mini-batch:

$$\text{Loss} = \frac{1}{8} \sum_{m=1}^8 \left[\sum_{s=1}^3 \left(- \sum_{k=1}^2 y_{\text{msk}} \cdot \log(p_{\text{msk}}) \right) \right]$$

Based on the mini-batch loss the weights in the tow weight matrices of layer 1 and layer 2 are updated.

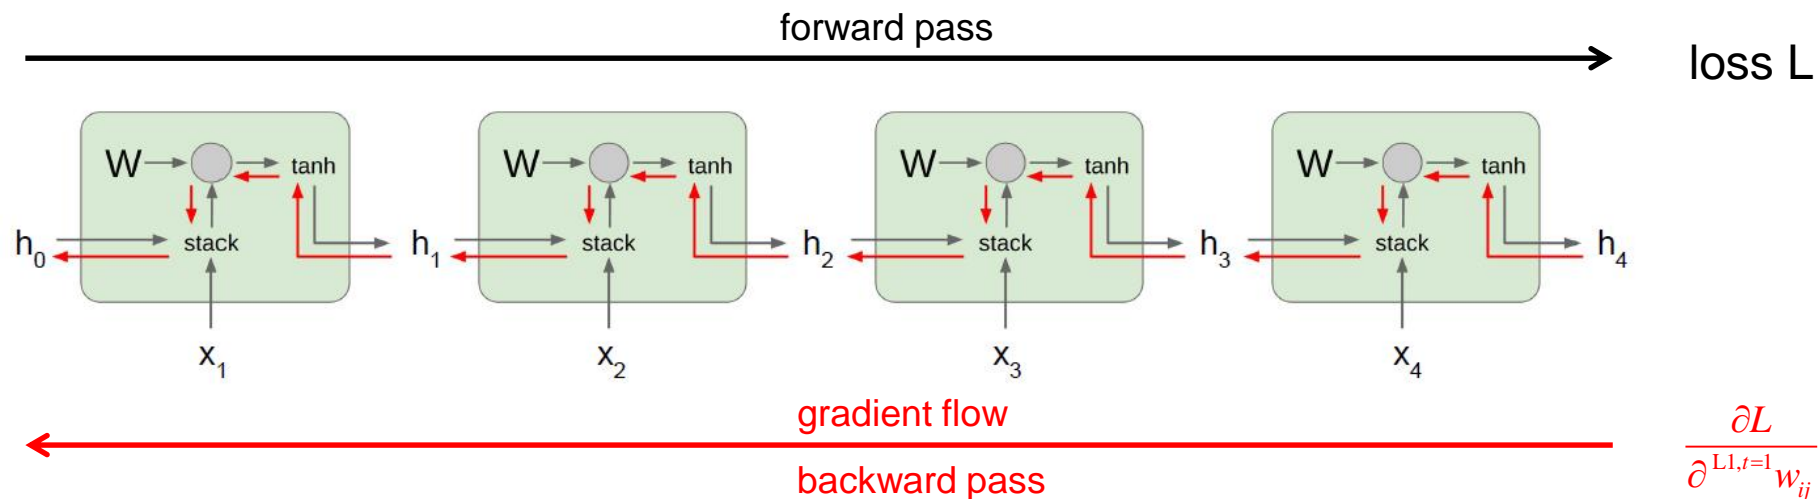
Recall: Design of a RNN “cell”



parameter?

$$(|h| + |x|) \cdot |h| + |h|$$

Backpropagation in RNNs: Gradient is multiplied at each time step with same factor: Gradient explosion/vanishing



Propagating the gradient of the cost function via chain rule to the first time point involves multiplying at each time step with \mathbf{W}^T (and the derivation of \tanh).

⇒ Vanishing gradient if we multiply at each time step with a number < 1

(more precisely we have only a number if W is a scalar, otherwise we need to look on the first singular value of \mathbf{W}^T)

⇒ Exploding gradient if we multiply at each time step with a number > 1

(more precisely we have only a number if W is a scalar, otherwise we need to look on the first singular value of \mathbf{W}^T)


Solution: gradient clipping (hack), or use better architecture like LSTM or GRU!

GRU and LSTM cells to avoid
vanishing/exploding gradients

Recall: ResNet

- use ResNet like architectures allowing for a gradient highway

(in CNN also batch-normalization and ReLU helped to train deep NN, but cannot naively transferred to recurrent NN)



ResNet basic design (VGG-style)

- add shortcut connections every two
- all 3x3 conv (almost)

152 layers:


Why does this train at all?

This deep architecture could still be trained, since the gradients can skip layers which diminish the gradient!


$$x_{l+1} = x_l + F(x_l)$$

Provide gradient highway also in recurrent NN: GRU, LSTM


CNN classic:



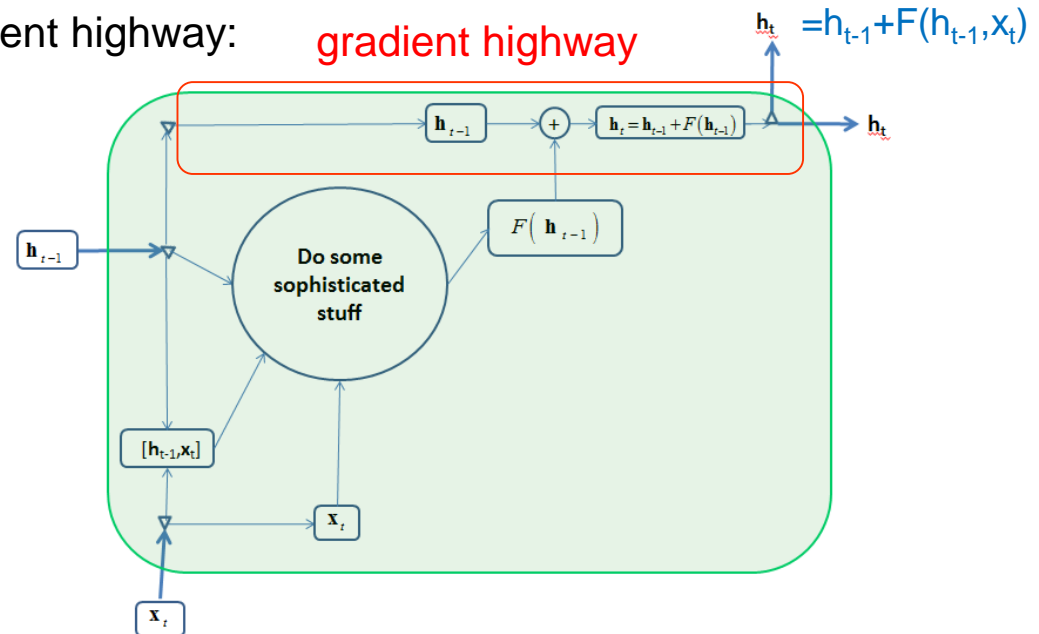
CNN with gradient highway:



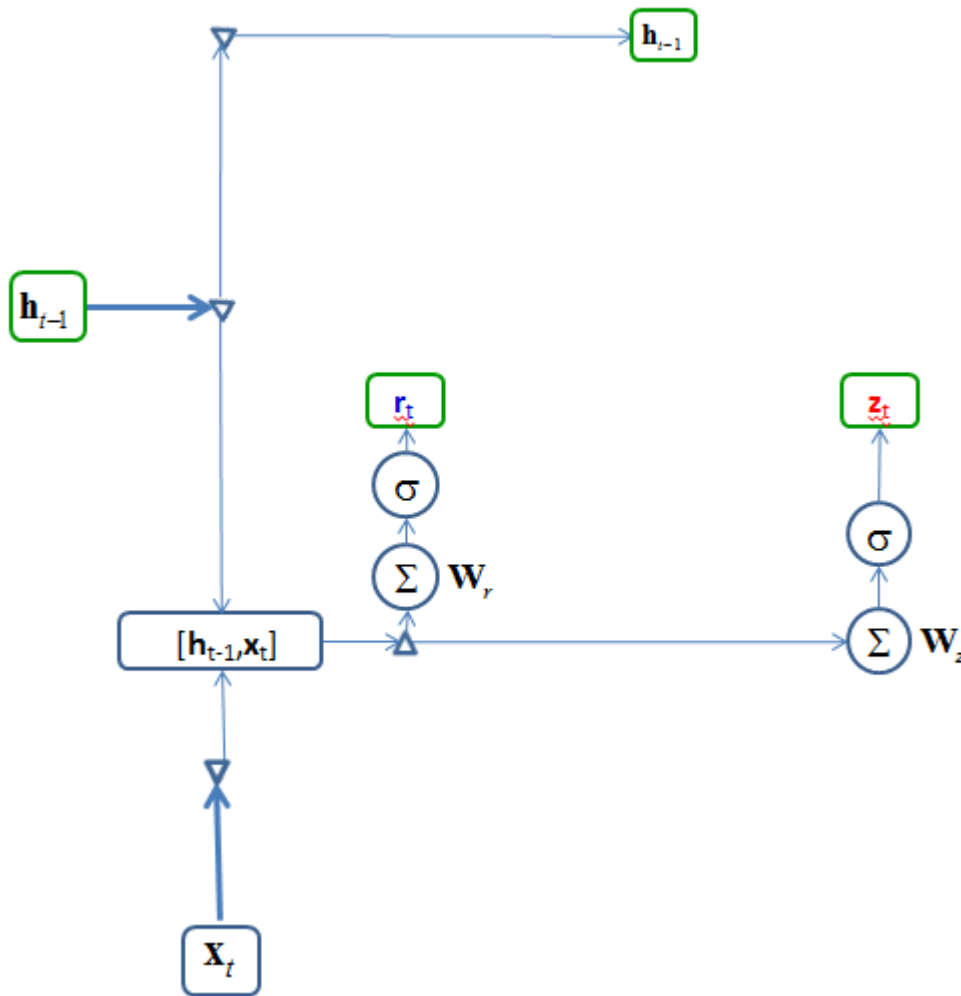
RNN classic:



RNN with gradient highway:



Towards Gated Recurrent Units (GRU)



$$\mathbf{r}_t = \text{gate}_{r,t} = \text{sigmoid}([\mathbf{h}_{t-1}, \mathbf{x}_t] \cdot \mathbf{W}_r + \mathbf{b}_r)$$

$$\mathbf{z}_t = \text{gate}_{\text{update}} = \text{sigmoid}([\mathbf{h}_{t-1}, \mathbf{x}_t] \cdot \mathbf{W}_z + \mathbf{b}_z)$$

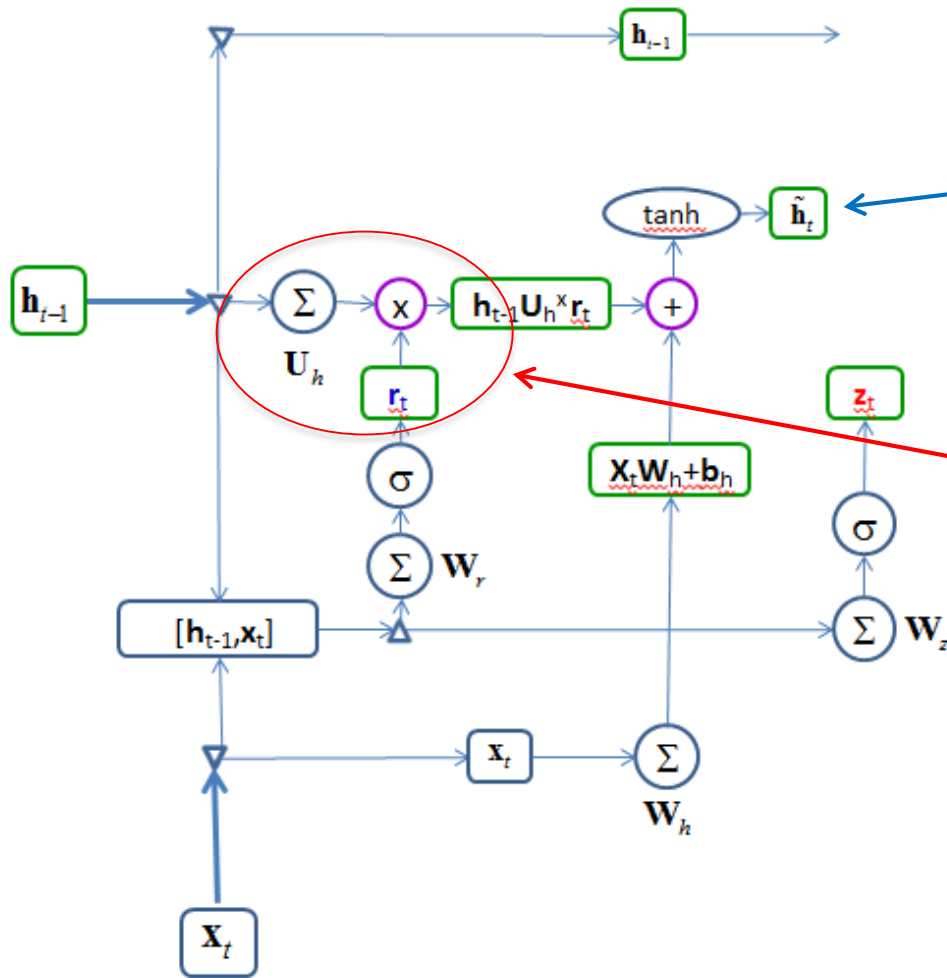
Idea:

The **relevant gate r** controls which part of the previous hidden state is relevant for making a prediction or should be dropped

The **updated gate z** controls how much information from the previous hidden layer h_{i-1} and the new input should be propagated to the current hidden layer h_i .

Remark: all internal vectors/tensors with green frame have same length/shape.

Towards Gated Recurrent Units (GRU)



The new proposed state \tilde{h} is:

$$\tilde{h}_t = \tanh(\mathbf{x}_t \cdot \mathbf{W}_h + \mathbf{b}_h + \mathbf{h}_{t-1} \mathbf{U}_h \otimes \mathbf{r}_t)$$

Here, we use the **relevant gate r** to control what part of \mathbf{h}_{t-1} we need to compute a new proposal.

Remark: all internal vectors with **green frame** have same length.

All ops within a **purple circle** are performed per **element-wise** on the ingoing vectors

The Gated Recurrent Unit (GRU)

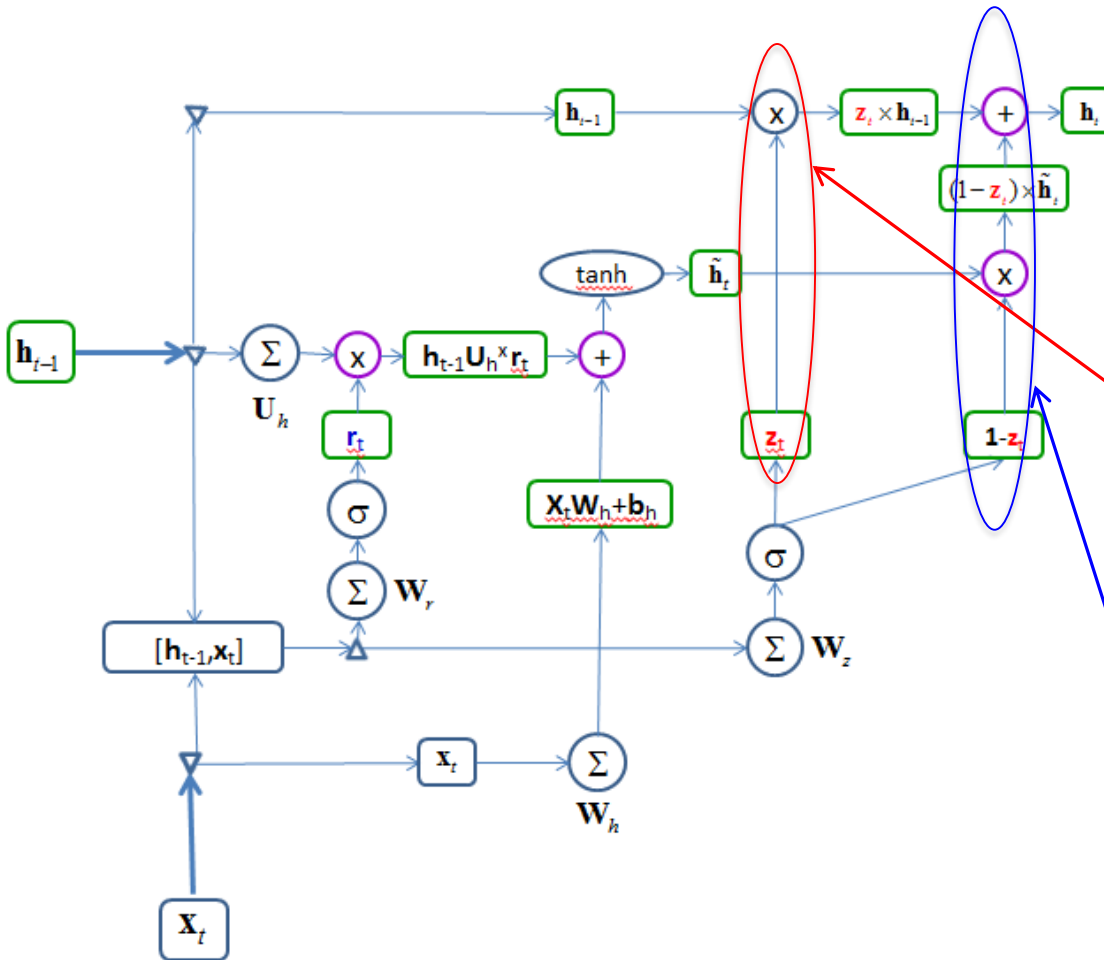
The new hidden state is:

$$\mathbf{h}_t = (\mathbf{1} - \mathbf{z}_t) \otimes \tilde{\mathbf{h}}_t \oplus \mathbf{z}_t \otimes \mathbf{h}_{t-1}$$

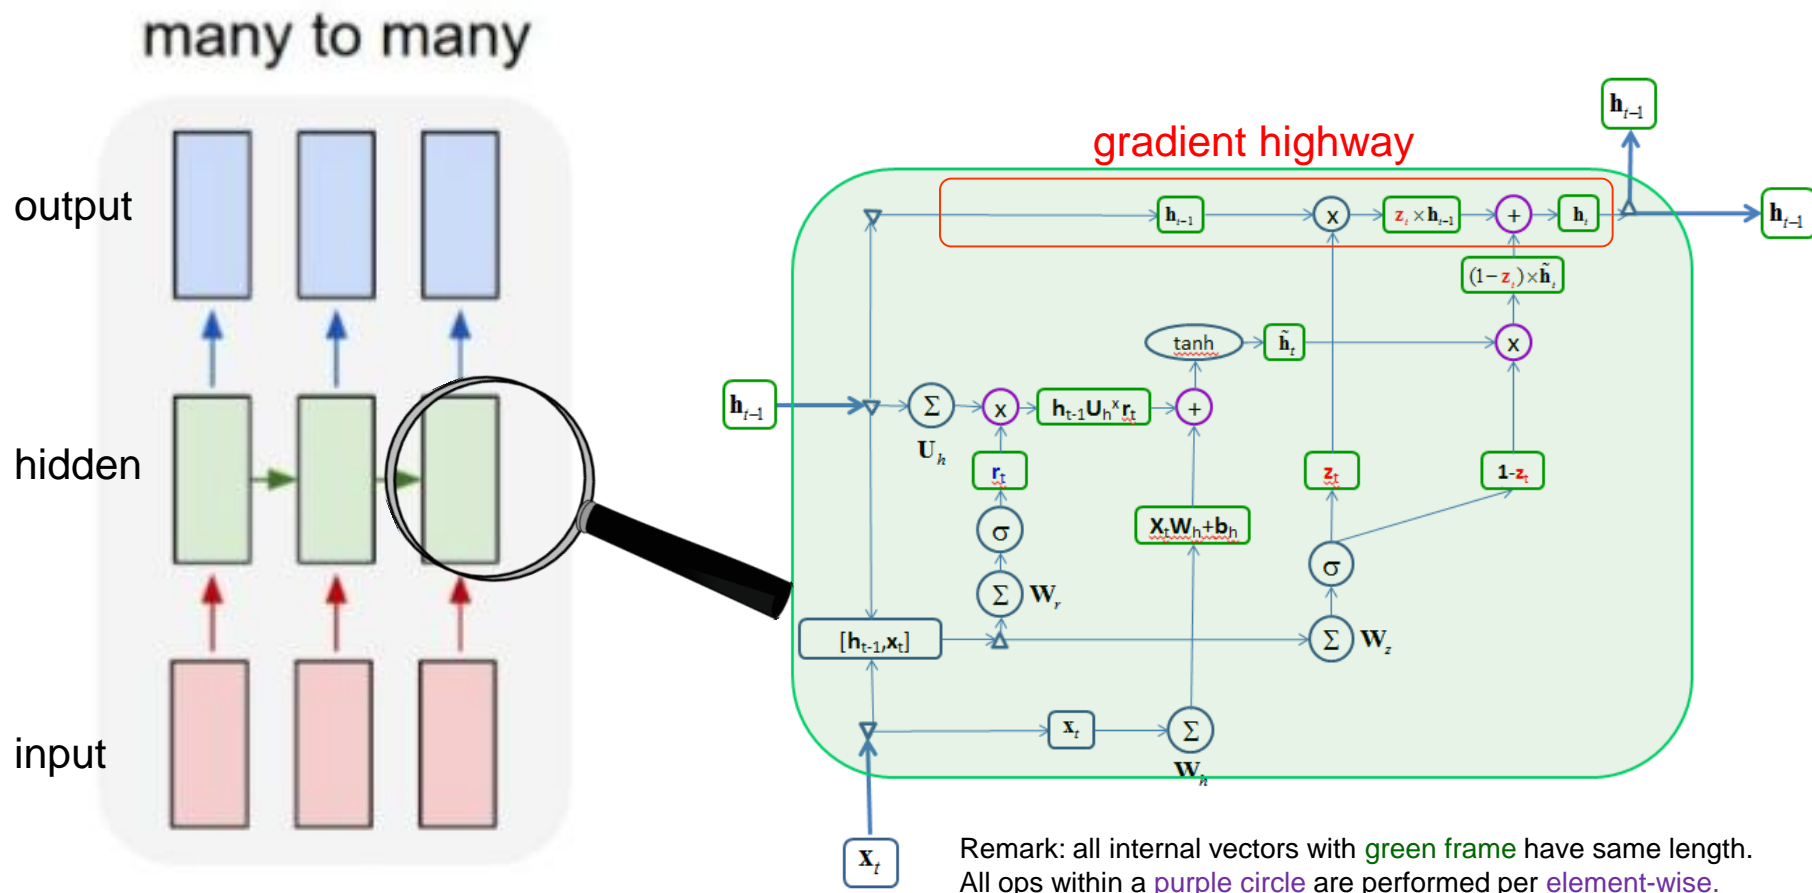
If all elements of \mathbf{z}_t are 1 then the hidden state stays unchanged.

The updated gate \mathbf{z}_t controls how much of the previous hidden state \mathbf{h}_{t-1} and the new input \mathbf{x}_t should be propagated to the current hidden state \mathbf{h}_t

The updated gate \mathbf{z}_t controls also how much information from proposed new state $\tilde{\mathbf{h}}$ is entering the new state



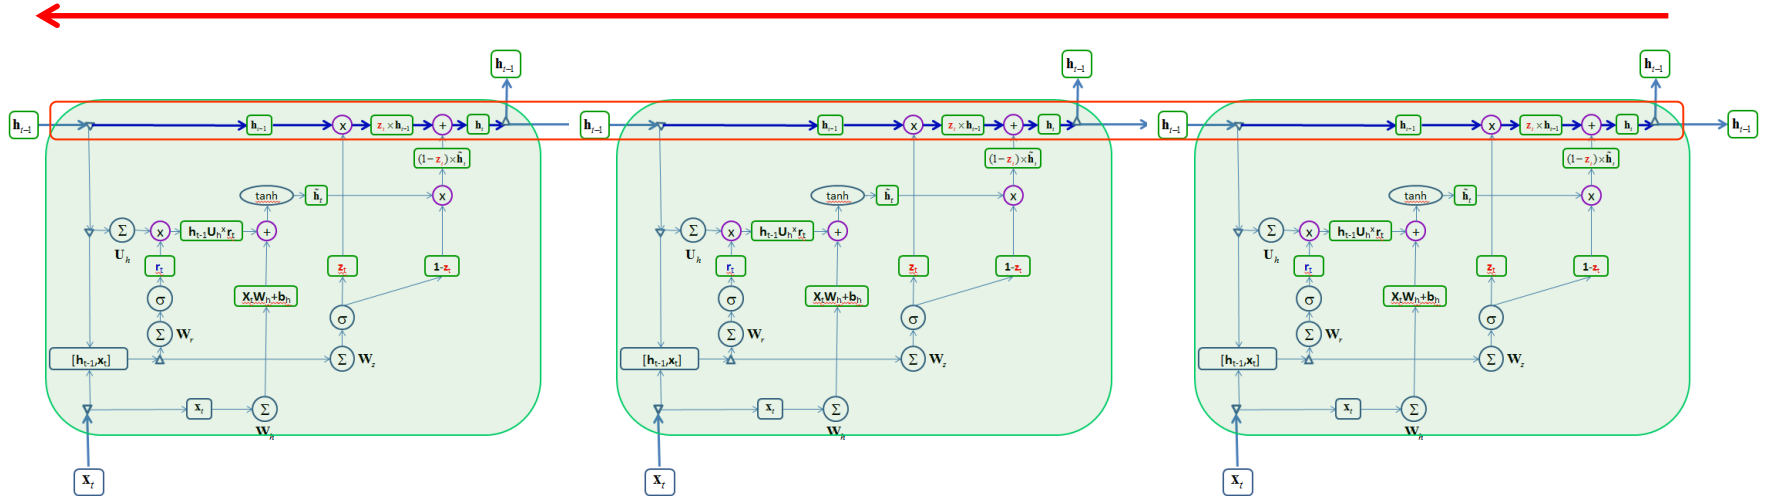
Solution via “highway allowing” architecture: GRU



The gradient **high-way** avoids **gradient vanishing**. The **GRU** also avoids **gradient explosion** since the element-wise operations on vector-elements that change over the time steps, avoids multiplying the gradients with the same number in each step.

The Gated Recurrent Unit (GRU): Gradient Flow

Uninterrupted gradient flow!



parameter?

Relevant gate: $\mathbf{r}_t = \sigma([\mathbf{h}_{t-1}, \mathbf{x}_t] \cdot \mathbf{W}_r + \mathbf{b}_r)$

Update gate: $\mathbf{z}_t = \sigma([\mathbf{h}_{t-1}, \mathbf{x}_t] \cdot \mathbf{W}_z + \mathbf{b}_z)$

Proposed hidden state: $\tilde{\mathbf{h}}_t = \tanh(\mathbf{x}_t \cdot \mathbf{W}_h + \mathbf{b}_h + \mathbf{h}_{t-1} \mathbf{U}_h \otimes \mathbf{r}_t)$

New hidden state is: $\mathbf{h}_t = (\mathbf{1} - \mathbf{z}_t) \otimes \tilde{\mathbf{h}}_t \oplus \mathbf{z}_t \otimes \mathbf{h}_{t-1}$

$$\begin{aligned} & (|h| + |x|) \cdot |h| + |h| \\ & + (|h| + |x|) \cdot |h| + |h| \\ & + (|x|) \cdot |h| + |h| + (|h|) \cdot |h| \end{aligned}$$

$$= 3 \cdot [(|h| + |x|) \cdot |h| + |h|]$$

A simplified variation, the Gated Recurrent Unit, or GRU, introduced by [Cho, et al. \(2014\)](#).


GRU in keras

```
from keras.models import Sequential
from keras import layers
from keras.optimizers import RMSprop

model = Sequential()
model.add(layers.GRU(32, input_shape=(None, float_data.shape[-1])))
model.add(layers.Dense(1))

model.compile(optimizer=RMSprop(), loss='mae')
history = model.fit_generator(train_gen,
                              steps_per_epoch=500,
                              epochs=20,
                              validation_data=val_gen,
                              validation_steps=val_steps)
```

length of internal state




All internal vectors within GRU green frame have same length.

All ops within a purple circle are performed per element-wise on the ingoing vectors

Long Short Term Memory (LSTM): Gradient Flow

LSTM has an additional cell state C for a “long term memory”.

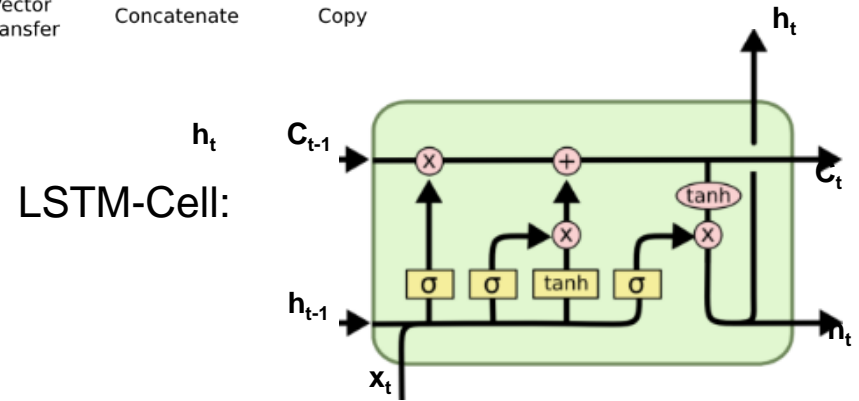


Uninterrupted gradient flow!



Long Short Term Memory networks (LSTM) were introduced by [Hochreiter & Schmidhuber \(1997\)](#).

Slide credit (modified): cs231 2017 stanford

Long Short Term Memory cell (LSTM) as GRU-extension



2 gates, 1 cell states (h)

Relevant gate: $\mathbf{r}_t = \sigma([\mathbf{h}_{t-1}, \mathbf{x}_t] \cdot \mathbf{W}_r + \mathbf{b}_r)$

Update gate: $\mathbf{z}_t = \sigma([\mathbf{h}_{t-1}, \mathbf{x}_t] \cdot \mathbf{W}_z + \mathbf{b}_z)$

Proposed hidden state: $\tilde{\mathbf{h}}_t = \tanh(\mathbf{x}_t \cdot \mathbf{W}_h + \mathbf{b}_h + \mathbf{h}_{t-1} \mathbf{U}_h \otimes \mathbf{r}_t)$

New hidden state is: $\mathbf{h}_t = (\mathbf{1} - \mathbf{z}_t) \otimes \tilde{\mathbf{h}}_t \oplus \mathbf{z}_t \otimes \mathbf{h}_{t-1}$

3 gates, 2 cell states (S:h, L:C)

Forget gate: $\mathbf{f}_t = \sigma([\mathbf{h}_{t-1}, \mathbf{x}_t] \cdot \mathbf{W}_f + \mathbf{b}_f)$

Input gate: $\mathbf{i}_t = \sigma([\mathbf{h}_{t-1}, \mathbf{x}_t] \cdot \mathbf{W}_i + \mathbf{b}_i)$

Output gate: $\mathbf{o}_t = \sigma([\mathbf{h}_{t-1}, \mathbf{x}_t] \cdot \mathbf{W}_o + \mathbf{b}_o)$

Proposed cell state: $\tilde{\mathbf{C}}_t = \tanh([\mathbf{h}_{t-1}, \mathbf{x}_t] \cdot \mathbf{W}_C + \mathbf{b}_C)$

New L cell state: $\mathbf{C}_t = \mathbf{f}_t \otimes \mathbf{C}_{t-1} \oplus \mathbf{i}_t \otimes \tilde{\mathbf{C}}_t$

New S hidden state: $\mathbf{h}_t = \mathbf{o}_t \otimes \tanh(\mathbf{C}_t)$

Long Short Term Memory (LSTM) in keras

```
from keras.layers import LSTM

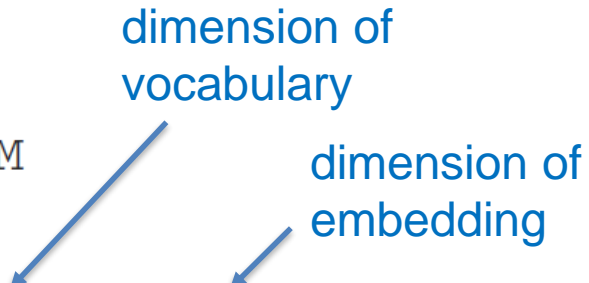
model = Sequential()
model.add(Embedding(max_features, 32))
model.add(LSTM(32))
model.add(Dense(1, activation='sigmoid'))

model.compile(optimizer='rmsprop',
              loss='binary_crossentropy',
              metrics=['acc'])

history = model.fit(input_train, y_train,
                    epochs=10,
                    batch_size=128,
                    validation_split=0.2)
```

dimension of vocabulary

dimension of embedding



Model zoo: many pretrained NN are out there

<https://modelzoo.co/>

Base pretrained models and datasets in pytorch (MNIST, SVHN, CIFAR10, CIFAR100, STL10, AlexNet, VGG16, VGG19, ResNet, Inception, SqueezeNet)

The Model Zoo website displays a grid of 48 model cards, each representing a different pre-trained neural network. The cards are organized into 4 columns and 12 rows. Several cards are circled in red, highlighting specific models:


- imagnet-vgg**: A VGG16 model trained on the ImageNet dataset, supporting PyTorch, TensorFlow, and Keras.
- PyTorch Image Classification with Kaggle Dogs vs Cats Dataset**: A VGG16 model trained on the Kaggle Dogs vs Cats dataset, supporting PyTorch, TensorFlow, and Keras.
- Bidirectional LSTM on the IMDB dataset**: A Bidirectional LSTM model trained on the IMDB dataset, supporting Keras.
- 1D CNN on the IMDB dataset**: A 1D CNN model trained on the IMDB dataset, supporting Keras.
- 1D CNN-LSTM on the IMDB dataset**: A 1D CNN-LSTM model trained on the IMDB dataset, supporting Keras.
- Using pre-trained word embeddings**: A model using pre-trained word embeddings, supporting Keras.
- Simple CNN on MNIST**: A Simple CNN model trained on the MNIST dataset, supporting Keras.
- Simple CNN with data augmentation**: A Simple CNN model trained on the CIFAR10 dataset with data augmentation, supporting Keras.

Other visible model cards include:

- Graphics code generating model using Processing**: A model for generating graphics code, supporting PyTorch and TensorFlow.
- chainer-Variational-AutoEncoder**: A Variational AutoEncoder (VAE) model, supporting Chainer and TensorFlow.
- Hierarchical Attention Network for Document Classification**: A Hierarchical Attention Network (HAN) model for document classification, supporting PyTorch and TensorFlow.
- improved-gan**: An improved Generative Adversarial Network (GAN) model, supporting Chainer and TensorFlow.
- Simple Generative Adversarial Networks**: A Simple GAN model, supporting PyTorch and TensorFlow.
- multilabel**: A multilabel classification model, supporting TensorFlow and Keras.
- VAE**: A Variational AutoEncoder (VAE) model, supporting Chainer and TensorFlow.
- SqueezeNet**: A SqueezeNet model, supporting TensorFlow and Keras.
- chainer_encoder_decoder**: A Chainer Encoder-Decoder model, supporting Chainer and TensorFlow.
- ADDA**: An Adversarial Discriminative Domain Adaptation (ADDA) model, supporting Chainer and TensorFlow.
- mxnet-audio**: An MXNet audio model, supporting MXNet and TensorFlow.
- MXSeq2Seq(Gluon:star**: An MXNet sequence-to-sequence model, supporting MXNet and TensorFlow.
- Inception v3**: An Inception v3 model, supporting Keras.
- Neural Style Transfer**: A Neural Style Transfer model, supporting Keras.
- Visualizing the filters learned by a CNN**: A model for visualizing CNN filters, supporting Keras.
- Deep dreams**: A Deep Dreams model, supporting Keras.
- Stateful LSTM**: A Stateful LSTM model, supporting Keras.
- Siamese network**: A Siamese network model, supporting Keras.
- DeconvNet**: A DeconvNet model for semantic segmentation, supporting TensorFlow and Keras.
- Pixel-wise Segmentation on VOC2012 Dataset using PyTorch**: A Pixel-wise Segmentation model trained on the VOC2012 dataset, supporting PyTorch and TensorFlow.
- generative-models**: A collection of generative models, supporting PyTorch and TensorFlow.
- V-Net**: A V-Net model for volumetric medical image segmentation, supporting PyTorch and TensorFlow.
- Evolution Strategies**: A model using Evolution Strategies for reinforcement learning, supporting PyTorch and TensorFlow.
- CycleGAN and Semi-Supervised GAN**: A CycleGAN and Semi-Supervised GAN model, supporting PyTorch and TensorFlow.
- Adversarial Generator-Encoder Network**: An Adversarial Generator-Encoder Network model, supporting PyTorch and TensorFlow.
- Recurrent Variational Autoencoder**: A Recurrent Variational Autoencoder (RVAE) model, supporting PyTorch and TensorFlow.
- AttGAN**: An AttGAN model for face attribute transfer, supporting TensorFlow and Keras.
- BEGAN in PyTorch**: A BEGAN model in PyTorch, supporting PyTorch and TensorFlow.
- Neural machine translation between the writings of Shakespeare and modern English using TensorFlow**: A Neural Machine Translation (NMT) model, supporting TensorFlow and Keras.
- img_classification_pk_pytorch**: A PyTorch image classification model, supporting PyTorch and TensorFlow.
- PNASNet.pytorch**: A PNASNet model in PyTorch, supporting PyTorch and TensorFlow.
- U-Net**: A U-Net model for brain tumor segmentation, supporting TensorFlow and Keras.
- CNN-LSTM-CTC**: A CNN-LSTM-CTC model for text recognition, supporting TensorFlow and Keras.

Can LSTM improve your conv1D series predictions?

- Work through the instructions in the second exercise in day 6 using https://github.com/tensorchiefs/dl_course_2018/blob/master/notebooks/12_LSTM_vs_1DConv.ipynb



Exercise: predictions with numeric time-series

a) Open the notebook [LSTM_vs_1DConv](#)

Look at the data generating process and train the first “1D Convolution without dilation rate” model and look at the predictions for the next 10 and 80 time steps.



What to you observe, did the model learn the data generating process?

How big is the receptive field in the last convolutional layer?

```
model_1Dconv = Sequential()
ks = 5
model_1Dconv.add(Convolution1D(filters=32, kernel_size=ks, padding='causal', input_shape=(128, 1)))
model_1Dconv.add(Convolution1D(filters=32, kernel_size=ks, padding='causal'))
model_1Dconv.add(Convolution1D(filters=32, kernel_size=ks, padding='causal'))
model_1Dconv.add(Convolution1D(filters=32, kernel_size=ks, padding='causal'))
model_1Dconv.add(Dense(1))
model_1Dconv.add(Lambda(slice, arguments={'slice_length':look_ahead}))

model_1Dconv.compile(optimizer='adam', loss='mean_squared_error')
model_1Dconv.summary()
```

$$\begin{aligned} \text{ReceptiveField} &= f_0 \cdot d_0 + \sum_i (f_i - 1) \cdot d_i \\ &= 5 \cdot 1 + 4 \cdot 1 + 4 \cdot 1 + 4 \cdot 1 = 17 \end{aligned}$$



Exercise: predictions with numeric time-series

a) Open the notebook [LSTM_vs_1DConv](#)

Look at the data generating process and train the first “1D Convolution without dilation rate” model and look at the predictions for the next 10 and 80 time steps.



What do you observe, did the model learn the data generating process?

How big is the receptive field in the last convolutional layer?

```
model_1Dconv_w_d = Sequential()
ks = 5
model_1Dconv_w_d.add(Convolution1D(filters=32, kernel_size=ks, padding='causal', dilation_rate=1, input_shape=(128, 1)))
model_1Dconv_w_d.add(Convolution1D(filters=32, kernel_size=ks, padding='causal', dilation_rate=2))
model_1Dconv_w_d.add(Convolution1D(filters=32, kernel_size=ks, padding='causal', dilation_rate=4))
model_1Dconv_w_d.add(Convolution1D(filters=32, kernel_size=ks, padding='causal', dilation_rate=8))
model_1Dconv_w_d.add(Dense(1))
model_1Dconv_w_d.add(Lambda(slice, arguments={'slice_length': look_ahead}))

model_1Dconv_w_d.compile(optimizer='adam', loss='mean_squared_error')
model_1Dconv_w_d.summary()
```

$$\begin{aligned} \text{ReceptiveField} &= f_0 \cdot d_0 + \sum_i (f_i - 1) \cdot d_i \\ &= 5 \cdot 1 + 4 \cdot 2 + 4 \cdot 4 + 4 \cdot 8 = 61 \end{aligned}$$



Exercise: predictions with numeric time-series

c) Now, Let's use a RNN for the same process. How good are the predictions for the 10 and 80 time steps?

What does the argument `return_sequences` mean?

How many weights do we need if we use a hidden state size of 12, check your calculations with the `model.summary()`.

Did the model learn the data generating process?

```
model_simple_RNN = Sequential()



model_simple_RNN.add(SimpleRNN(12,return_sequences=True,input_shape=(128,1)))
model_simple_RNN.add(Dense(1))
model_simple_RNN.add(Lambda(slice, arguments={'slice_length':look_ahead}))

model_simple_RNN.summary()
model_simple_RNN.compile(optimizer='adam', loss='mean_squared_error')
```

Layer (type)	Output Shape	Param #
=====	=====	=====
simple_rnn_1 (SimpleRNN)	(None, 128, 12)	168
dense_3 (Dense)	(None, 128, 1)	13
lambda_3 (Lambda)	(None, 10, 1)	0
=====	=====	=====
Total params: 181		
Trainable params: 181		
Non-trainable params: 0		

Von (12+1) auf 12: $(12+1)*12+12=168$

Von (12) auf 1: $(12)*1+1=13$



Exercise: predictions with numeric time-series

d) Let's replace the RNN cell with a more complex LSTM cell, in keras LSTM.
Use the same size of units and the same architecture.
How many weights do we need now?



```
: model_LSTM = Sequential()

model_LSTM.add(LSTM(12,return_sequences=True,input_shape=(128,1)))
model_LSTM.add(Dense(1))
model_LSTM.add(Lambda(slice, arguments={'slice_length':look_ahead}))

model_LSTM.summary()
model_LSTM.compile(optimizer='adam', loss='mean_squared_error')
```

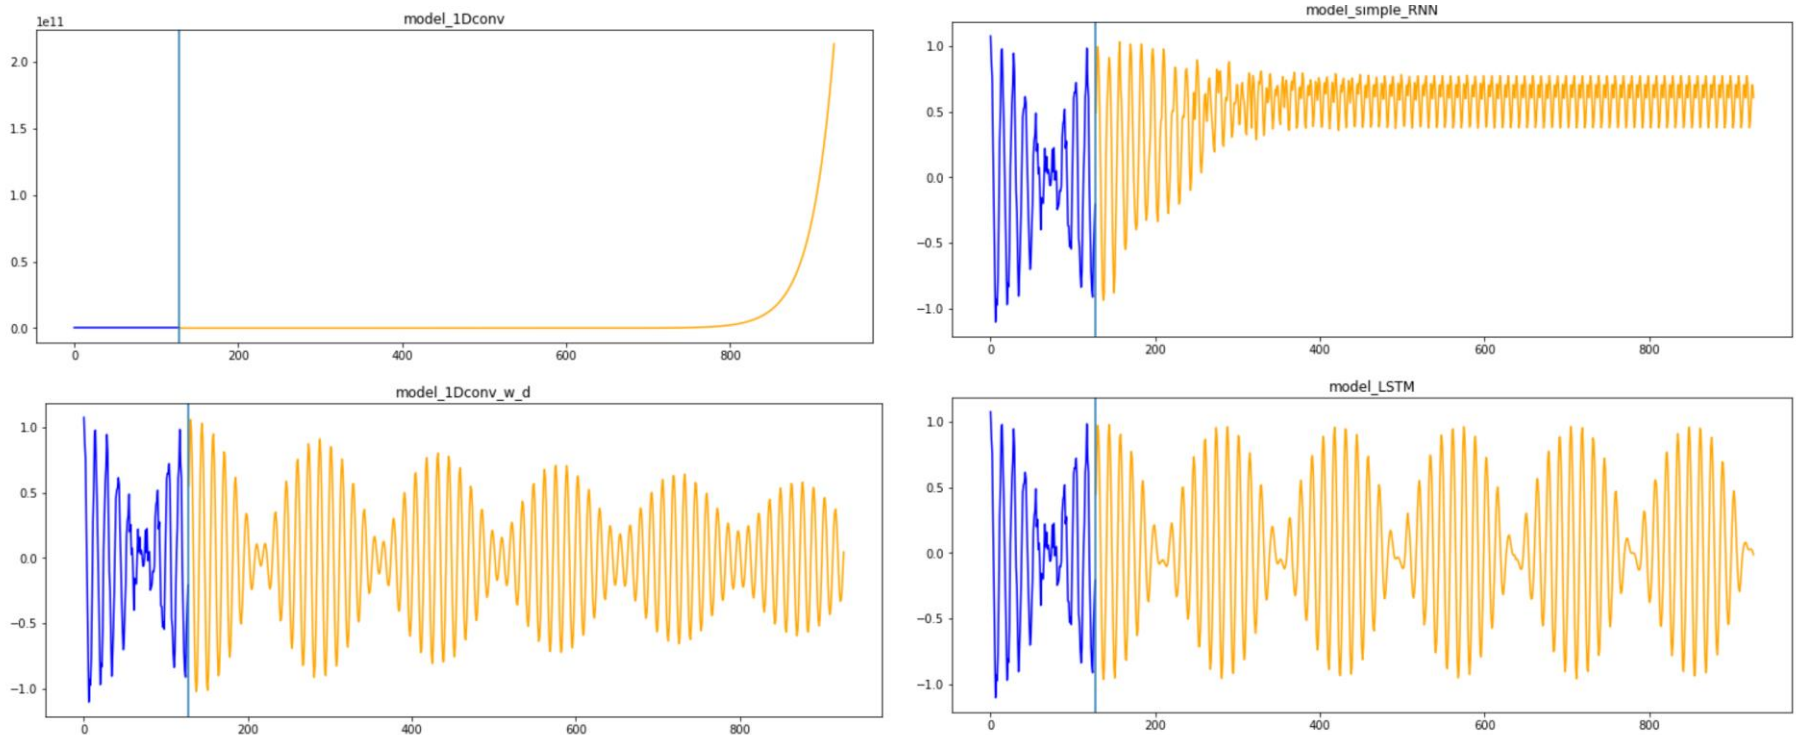
Layer (type)	Output Shape	Param #
lstm_1 (LSTM)	(None, 128, 12)	672
dense_4 (Dense)	(None, 128, 1)	13
lambda_4 (Lambda)	(None, 10, 1)	0
Total params: 685		
Trainable params: 685		
Non-trainable params: 0		

We have 4 –times from (h,x) to h
Here (12+1) to 12:
 $4 * ((12) * 13 + 12) = 672$



Exercise: predictions with numeric time-series

e) Compare the 4 models for very long predictions (800 timesteps). What do you observe? What could we do to improve the RNN and the Conv1D with dilation_rate?



To improve the performance, we would try to enlarge the memory and allow for more flexible models:

1D-CNN: stack more dilated layers to get a receptive field that sees whole input
RNN, LSTM: enlarge the dimension of the hidden state and stack more layers.

Time for your project